



MONASH University

Information Technology

FIT3143 - LECTURE WEEK 4

INTER-PROCESS COMMUNICATIONS & REMOTE
PROCEDURE CALL

algorithm distributed systems **database**
systems **computation** knowledge ma
design e-business **model** data mining **int**
distributed systems **database** software
computation knowledge management **an**

Overview


1. Inter process communication (IPC)
2. Synchronous and asynchronous communication
3. Remote procedure call

Learning outcome(s) related to this topic


- Compare and contrast different parallel computing architectures, algorithms and communication schemes using research-based knowledge and methods (LO2)

Inter-process Communication (IPC)

- IPC basically requires information sharing among two or more processes. Two basic methods-
 - Original sharing (Shared data approach)
 - Copy Sharing (Message passing approach)



Shared data approach



Message passing approach

Message Passing System

- A *message-passing system* is a sub-system of a distributed system that provides a set of message-based IPC protocols and does so by shielding the details of complex network protocols and multiple heterogeneous platforms from programmers.
- It enables processes to communicate by exchanging messages
- Allows programs to be written by using simple communication primitives, such as *send* and *receive*.
- Serves as a suitable infrastructure for building other higher level IPC systems, such as RPC (Remote Procedure Call) and DSM (Distributed Shared Memory).

Desirable features of a good message passing system

- Simplicity
- Uniform semantics
 - Same primitives for local and remote communication
- Efficiency
 - Reduce the number of message as far as possible
 - Some optimization normally adopted for efficiency include-
 - Avoiding the cost of establishing and terminating connections between the same pair of processes for each and every message exchange between them
 - Minimizing cost of maintaining connections
 - Piggybacking of acknowledgment of previous message with the next message.

Desirable features of a good message passing system

- Reliability
 - Lost and duplicate message handling
- Correctness
 - Atomicity
 - Ordered delivery
 - Survivability
- Flexibility
 - Can drop one or more correctness properties

Desirable features of a good message passing system

- Security
 - Authentication of sender and receiver
 - Encryption of messages
- Portability
 - Message passing system should itself be portable
 - The application written by using primitives of the IPC protocol should be portable.

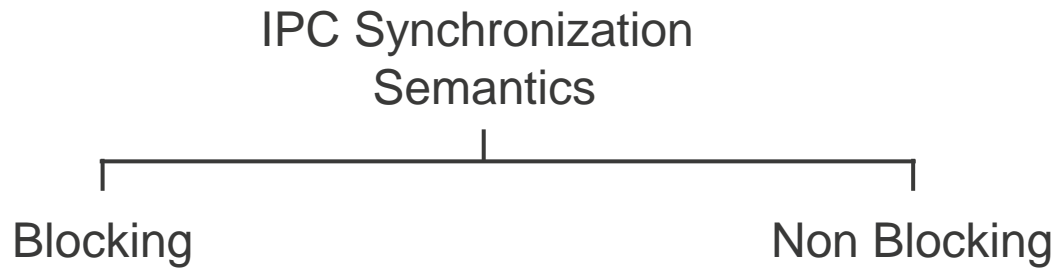
Issues in IPC by message passing

- A typical message structure
 - Header
 - Addresses
 - ✓ Sender address
 - ✓ Receiver address
 - Sequence number
 - Structural information
 - ✓ Type
 - ✓ Number of bytes
 - Message

Issues need to be considered for designing an IPC protocol

- Some of the main issues will be,
- Identity related
 - Who is the sender?
 - Who is the receiver?
- Network Topology related
 - 1 receiver or many?
- Flow control related
 - Guaranteed by the receiver?
 - Sender should wait for reply?
- Error control and channel management
 - Node crash.....what to do?
 - Receiver not ready....what to do?
 - Several outstanding messages for the receiver

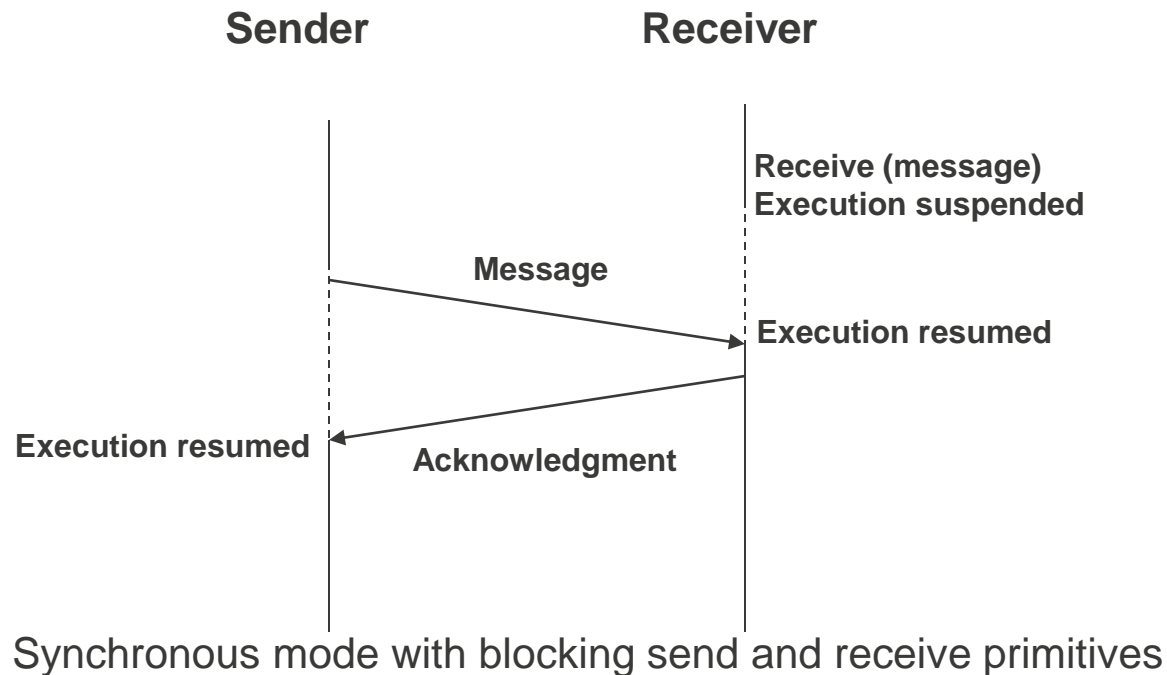
Synchronization in IPC



- Send primitive
 - Blocking
 - Non blocking
- Receive primitive
 - Blocking
 - Non Blocking
 - ✓ **Polling**
 - ✓ **interrupt**

Synchronous & Asynchronous Communication

- When both the send and receive primitives of a communication between two processes use blocking semantics, the communication is said to be Synchronous; otherwise it is asynchronous.



Synchronous vs Asynchronous Communication

- **Synchronous**

- Simple and easy to implement
- Contributes to reliability
- No backward error recovery needed

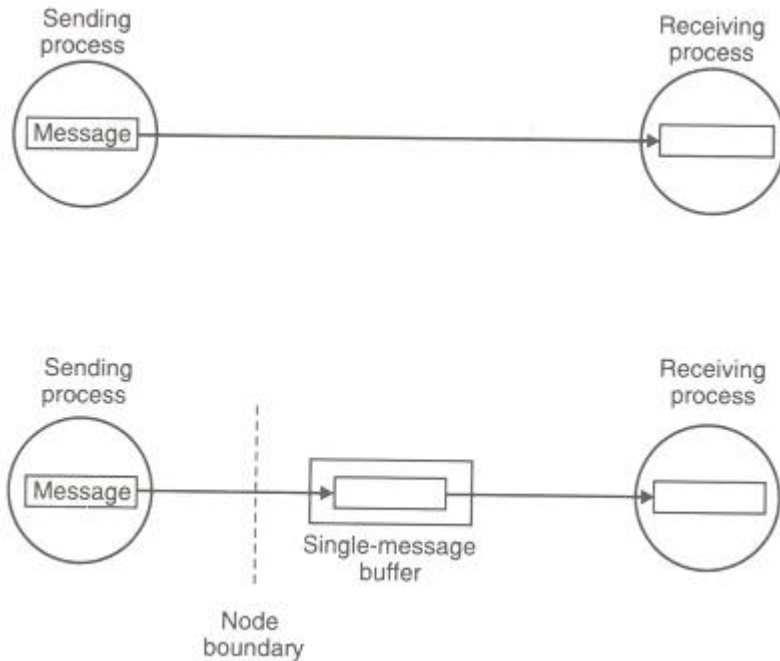
- **Asynchronous**

- High concurrency
- More flexible than synchronous
- Lower deadlock risk than in synchronous communication (but beware)


Buffering

- Synchronous systems
 - Null buffer
 - Single message buffer
- Asynchronous systems
 - Unbounded capacity buffer
 - Finite message (multiple message buffer)

Buffering



Synchronous System

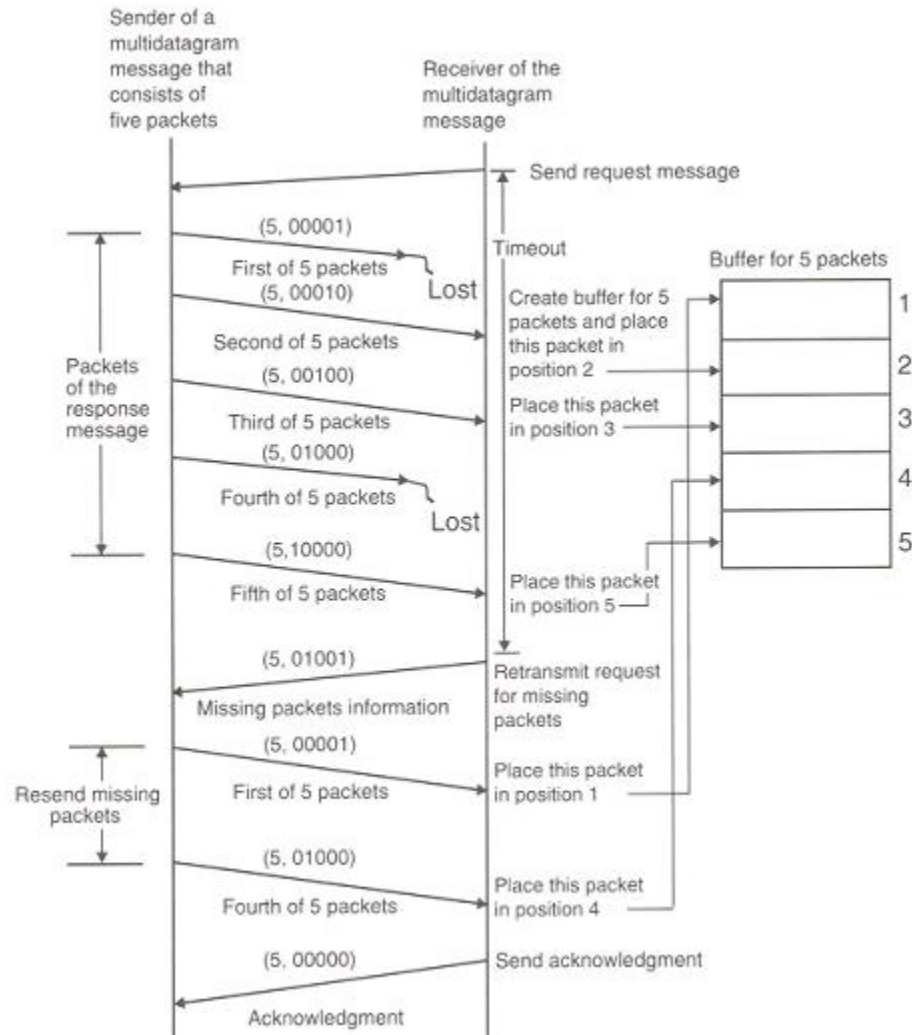


Asynchronous System

Multi-datagram Messages

- Almost all networks have an upper bound on the size of data that can be transmitted at a time. This is known as MTU (maximum transfer unit).
- Thus a message whose size is greater than MTU has to be fragmented- each fragment is sent in a packet. These packets are known as datagram.
- Thus messages may be single-datagram messages or multi-datagram messages
- Assembling and disassembling is the responsibility of message passing system.

Using Bitmap for multi datagrams



Encoding Decoding

Encoding/Decoding is needed if

- Sender and receiver have different architecture

Even for Homogeneous Encoding/Decoding is needed for

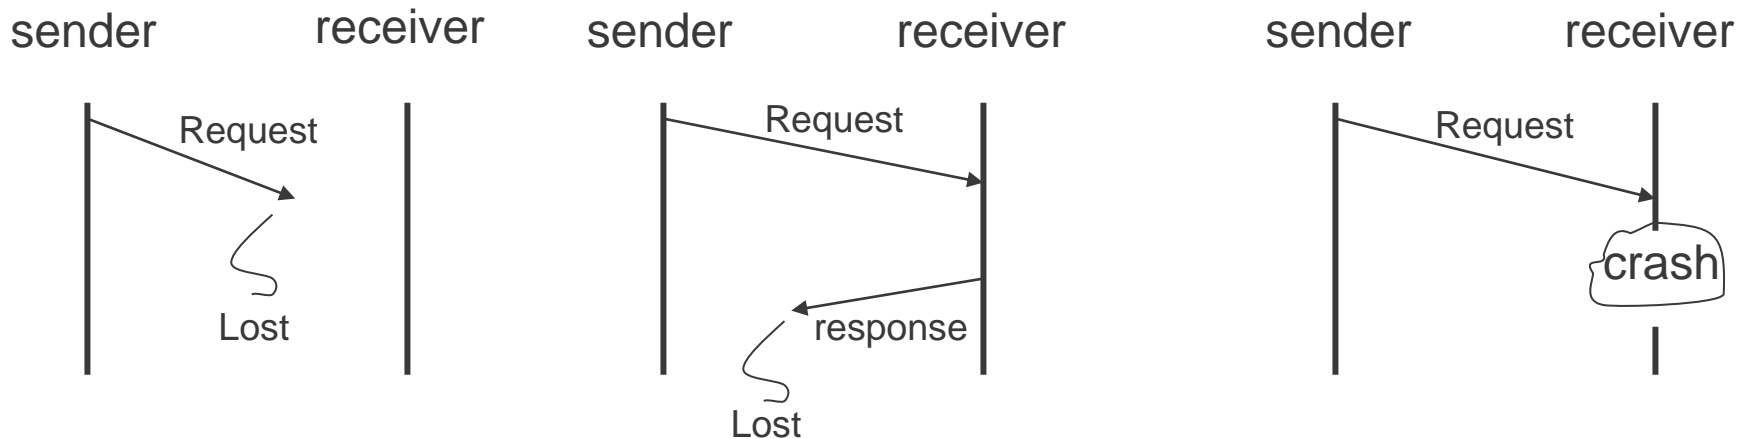
- ✓ Using an absolute pointer
- ✓ To know which object is stored in where and how much storage it requires

Process Addressing


- Explicit Addressing
 - Send(process_ID, message)
 - Receive(process_ID, message)
- Implicit Addressing
 - Send_any(service_ID, message)
 - Receive_any(process_ID, message)

Failure Handling

- Failure classification
 1. Loss of request message
 2. Loss of response message
 3. Unsuccessful execution of the request




Reliable IPC protocol




---- Blocked state
—— Executing state

4-message protocol



---- Blocked state
—— Executing state

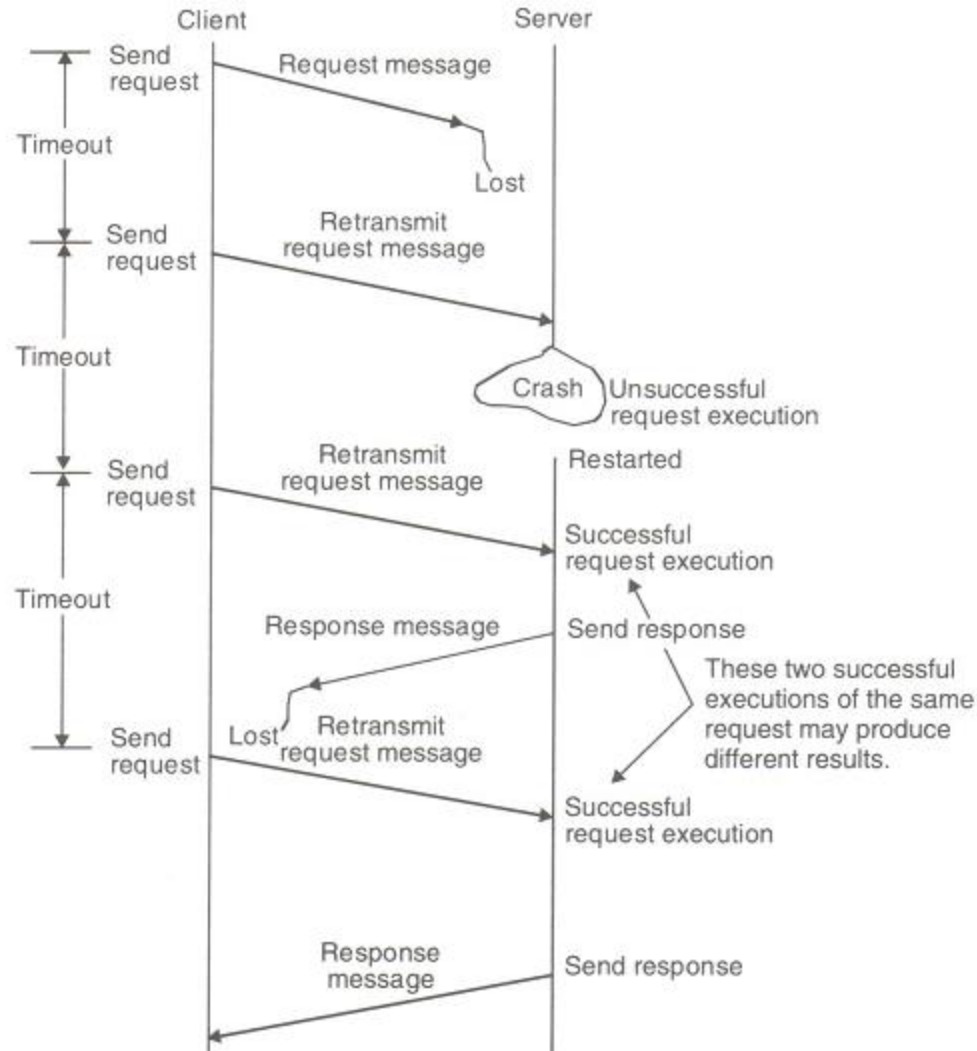
3-message protocol



---- Blocked state
—— Executing state

2-message protocol

An Example of Fault Tolerant System



Idempotency

- An idempotent function will return the same result given the same input, even when executed multiple times. For example, an idempotent function is illustrated in Algorithm 1; a nonidempotent function is illustrated in Algorithm 2.

Algorithm 1 An example of an idempotent function.

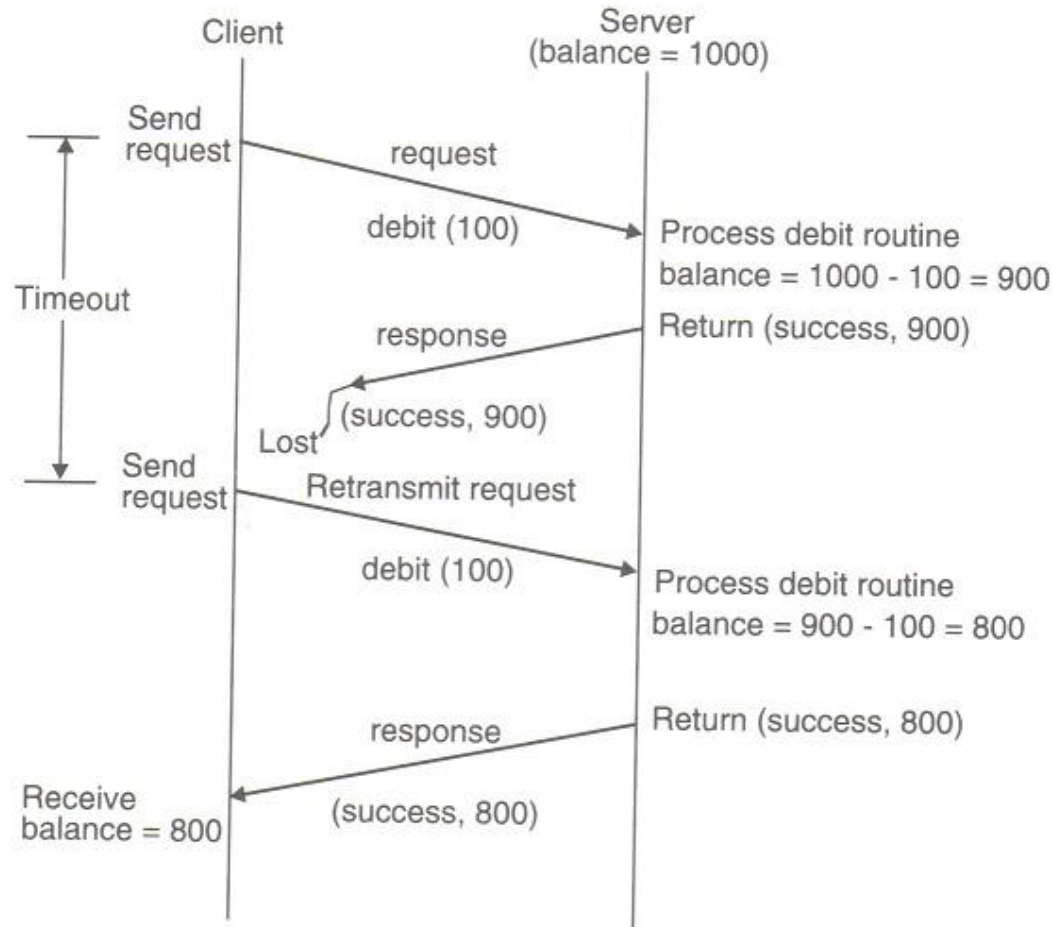
```
function GETSQRT(n)
    return SQRT(n)
end function
```

Algorithm 2 An example of a non-idempotent function. This function will return different values when given the same input. An illustration of the flow can be seen in [Figure 2.7](#).^{*}

```
function DEBIT(amount)
    if balance > amount then
        balance  $\leftarrow$  balance - amount
        return ("success", balance)
    else
        return ("failure", balance)
    end if
end function
```

**Note: Figure 2.7 refers to the following slide*

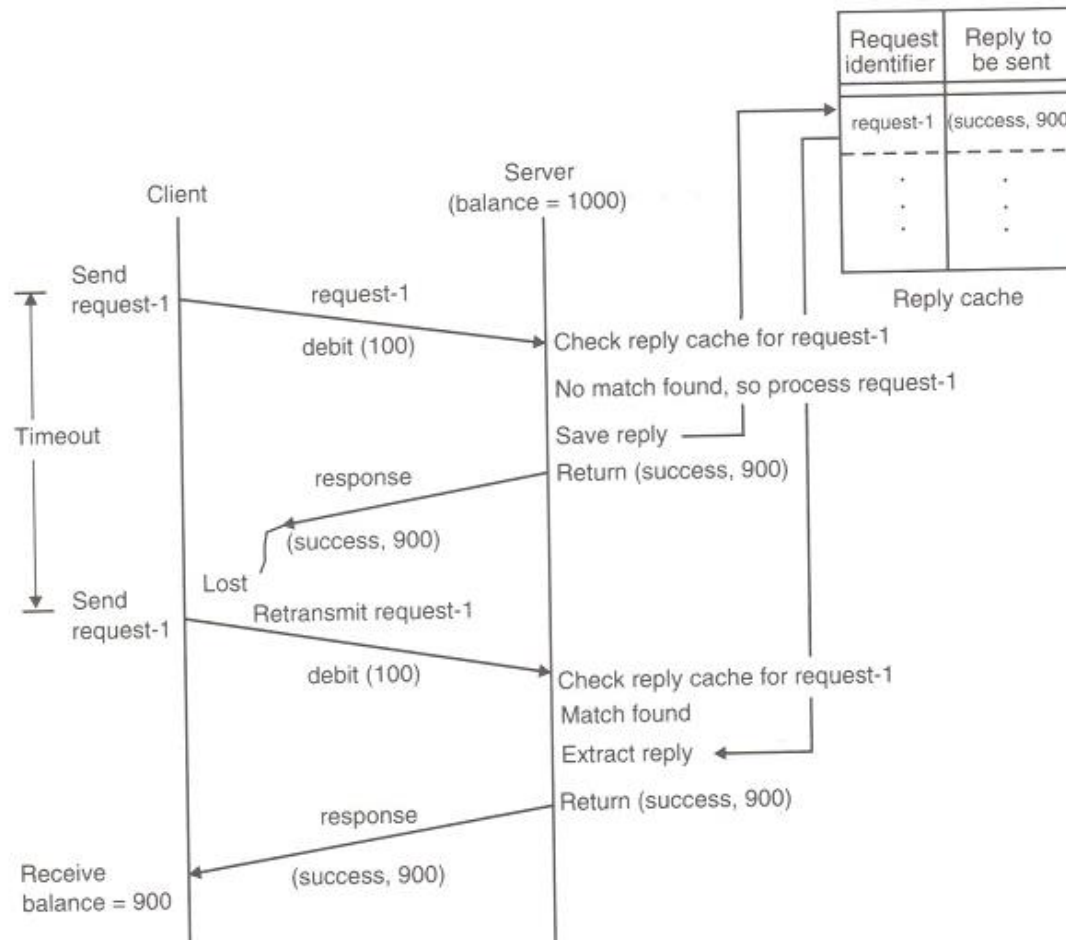
A Non-Idempotent Routine



Implementation of Idempotency

- How to implement Idempotency?
 - Adding sequence number with the request message
 - Introduction of 'Reply cache'

Implementation of Idempotency



Group Communication


- Group communication may be
 - One to many
 - Many to one
 - Many to many
- One to many
 - Group management
 - Group addressing
 - Buffered and unbuffered multicast
 - Send-to-all and Bulletin-Board semantics
 - Flexible reliability in multicast communication
 - Atomic multicast

Many to Many communication

- The issues related to one-to-many and many-to-one communications also applies here
- In addition, *ordered message delivery* is an important issue. This is trivial in one-to-many or many-to-one communications.
- For example, two server processes are maintaining a single salary database. Two client processes send updates for a salary record. What happen if they reach in different order? (will sequencing of messages help in this case?)

Semantics for ordered delivery in many-to-many comm.


- No-ordering



No ordering constraint

Semantics for ordered delivery in many-to-many comm.

- Absolute ordering
 - All messages are delivered to all receiver processes in the exact order in which they were sent.
 - Using global timestamp as message identifiers with sliding window protocol



Absolute ordering semantic


Semantics for ordered delivery in many-to-many comm.

- Consistent ordering
 - All messages are delivered to all receivers in the same order. However, this order may be different from the order in which messages were sent.

Implementation

- Make the communication appear as a combination of many-to-one and one-to-many communication [Chang and Maxemchuk]
- Kernels of sending machines send messages to a single receiver (known as *sequencer*) that assigns a sequence number to each message and then multicast it.
- Subject to single point of failure and hence has poor reliability.
- A distributed algorithm - *ABCAST* in *ISIS* system [Birman and Renesse] (self study)

Semantics for ordered delivery in many-to-many comm.



Consistent ordering semantic


Semantics for ordered delivery in many-to-many comm.

- Causal ordering
 - If the event of sending one message is causally related to the event of sending another message, the two messages are delivered to all receivers in correct order.
 - Two message sending events are said to be causally related if they are corelated by the *happened-before* relation.
- [The expression $a \rightarrow b$ is read “ a happens before b ” and means that all processes agree that first event a occurs, then afterward, event b occurs. The happens-before relation can be observed directly in two situations:
 1. If a and b are events in the same process, and a occurs before b , then $a \rightarrow b$ is true.
 2. If a is the event of a message being sent by one process, and b is the event of the message being received by another process, then $a \rightarrow b$ is also true.

Happens-before is a transitive relation, so if $a \rightarrow b$ and $b \rightarrow c$, then $a \rightarrow c$.]

Semantics for ordered delivery in many-to-many comm.

- One example of implementing Causal consistency is *CBCAST* in *ISIS* system [Birman et al].



Causal ordering semantic

Remote Procedure Call (RPC)

- The IPC part of a distributed application can be adequately and efficiently handled by using an IPC protocol based on message passing system.
- However, an independently developed IPC protocol is tailored specifically to one application and does not provide a foundation on which to build a variety of distributed applications.
- Therefore, a need was felt for a general IPC protocol that can be used for designing several distributed applications.
- The RPC facility emerged out of this need.


Remote Procedure Call

- While the RPC is not the universal panacea for all types of distributed applications but for a fairly large number of distributed applications.
- The RPC has become a widely accepted IPC mechanism in DS. Its features –
 - Simple call syntax.
 - Familiar semantics.
 - Specification of a well defined interface.
 - Ease of use.
 - Generality. “In single-machine computations procedure calls are often the most important mechanism for communication between the parts of the algorithm” [Birrell and Nelson].
 - Its efficiency


RPC model

- RPC model is similar to “Procedure call” model.
- Procedure call is same as function call or subroutine call
- Local Procedure Call - The caller and the callee are within a single process on a given host.
- Remote Procedure Call (RPC) - A process on the local system invokes a procedure on a remote system. The reason we call this a “procedure call” is because the intent is to make it appear to the programmer that a local procedure call is taking place.

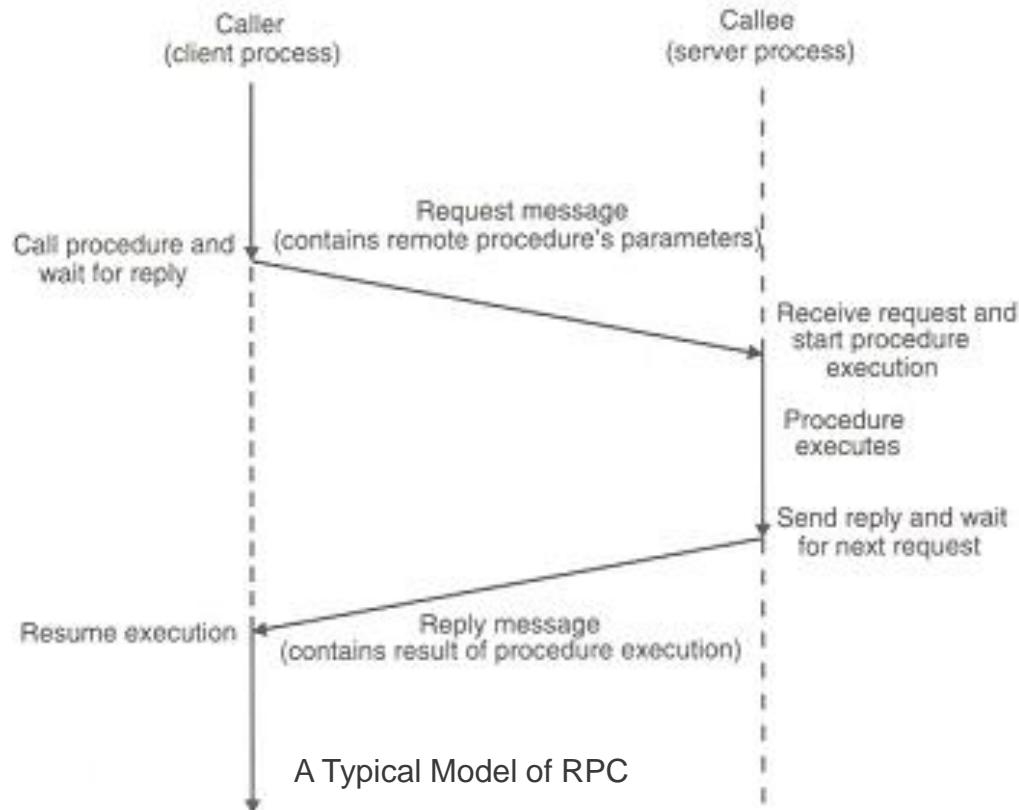
Local and Remote Procedure Call



Local Procedure Call



Remote Procedure Call

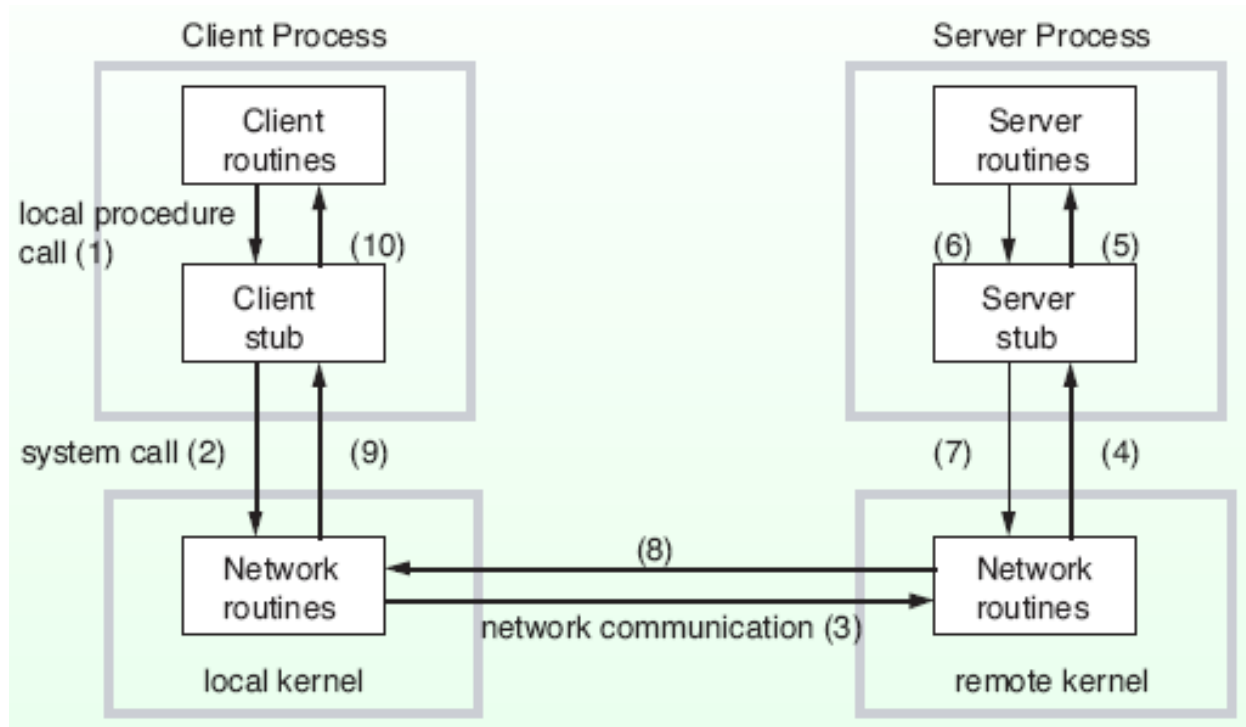


A Typical Model of RPC

Implementing RPC mechanism

- To achieve the goal of *semantic transparency*, the implementation of an RPC mechanism is based on the concept of *stubs*
- *Stubs* provide a perfectly normal(local) procedure call abstraction
- To hide the distance and functional details of underlying network, an RPC communication package (known as *RPCRuntime*) is used.
- Thus, RPC implementation involves five elements-
 - The client
 - The client stub
 - The RPCRuntime
 - The server stub
 - The server

RPC in Detail



Stubs


- Client and server stubs are generated from interface definition of server routines by development tools.
- Interface definition is similar to class definition in C++ and Java.

Parameter Passing Mechanisms

- When a procedure is called, parameters are passed to the procedure as the arguments. There are three methods to pass the parameters.
 - call-by-value
 - call-by-reference
 - call-by-copy/restore

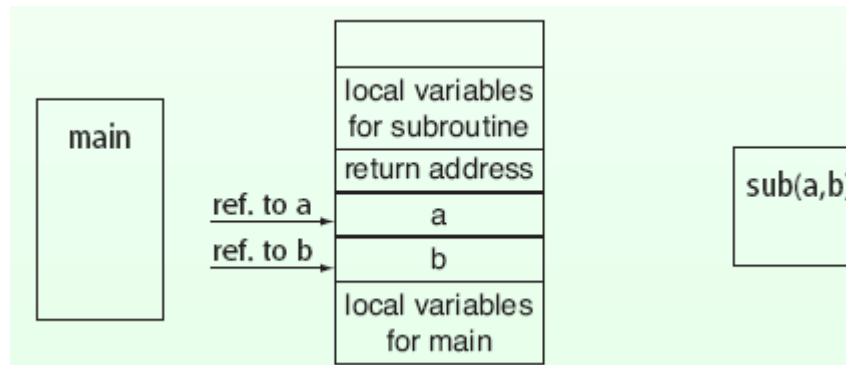
Call by value

- The values of the arguments are copied to the stack and passed to the procedure.
- The called procedure may modify these, but the modifications do not affect the original value at the calling side.




Call-by-Reference

- The memory addresses of the variables corresponding to the arguments are put into the stack and passed to the procedure.
- Since these are memory addresses, the original values at the calling side are changed if modified by the called procedure.



Call-by-Copy/Restore

- The values of the arguments are copied to the stack and passed to the procedure.
- When the processing of the procedure completes, the values are copied back to the original values at the calling side.
- If parameter values are changed in the subprogram, the values in the calling program are also affected.




Parameter Passing in RPC

- Which parameter passing mechanisms are possible?
 - It is possible to implement all of the three mechanisms if you wish. Usually call-by-value and call-by-copy/restore are used.
 - Call-by-reference is difficult to implement. All data which may be referenced must be copied to the remote host and the reference to the copied data is used.
- Do we need to convert the values of the arguments into a standard format to transmit over the network?


Parameter Passing in RPC

- Reasons to convert the values of the arguments into a standard format to transmit over the network
 - Different machines use different character codes. E.g., IBM main frames use EBCDIC, while PCs use ASCII.
 - Representation of numbers may differ from machine to machine.
 - Big endian and little endian




Parameter Passing in RPC

- If a standard format is not used, two message conversions are necessary.
 - If format information is attached to the message, only one conversion at the receiver will suffice.
 - However, the receiver must be able to handle many different formats.


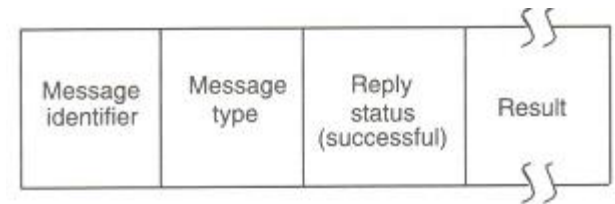


RPC Messages

- Generally, two types of messages
 - Call messages
 - Reply messages



A typical RPC Call message format

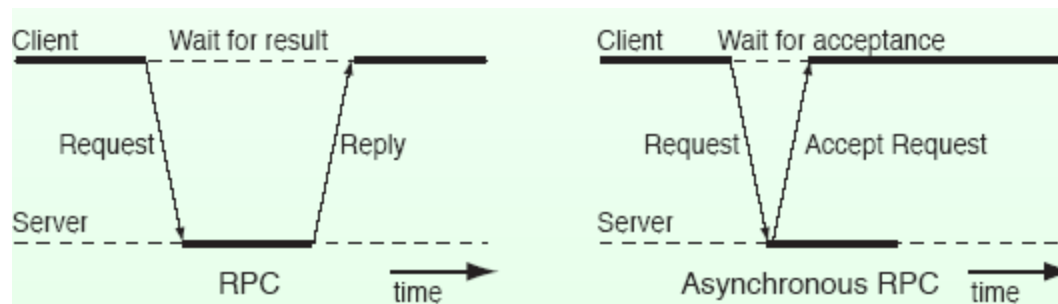


A typical RPC Reply message format (successful and unsuccessful)

Variations of RPC

Asynchronous RPC


- RPC (When a client requests a remote procedure, the client wait until a reply comes back in RPC).
- If no result is to be returned, unnecessary wait time overhead.
- In asynchronous RPC, the server immediately sends accept message when it receives a request.



Call-Back RPC

One-way RPC

- In one-way RPC, the client immediately continues after sending the request to the server.



Some special types of RPC

- Callback RPC
- Broadcast RPC
- Batch-mode RPC
- Lightweight RPC

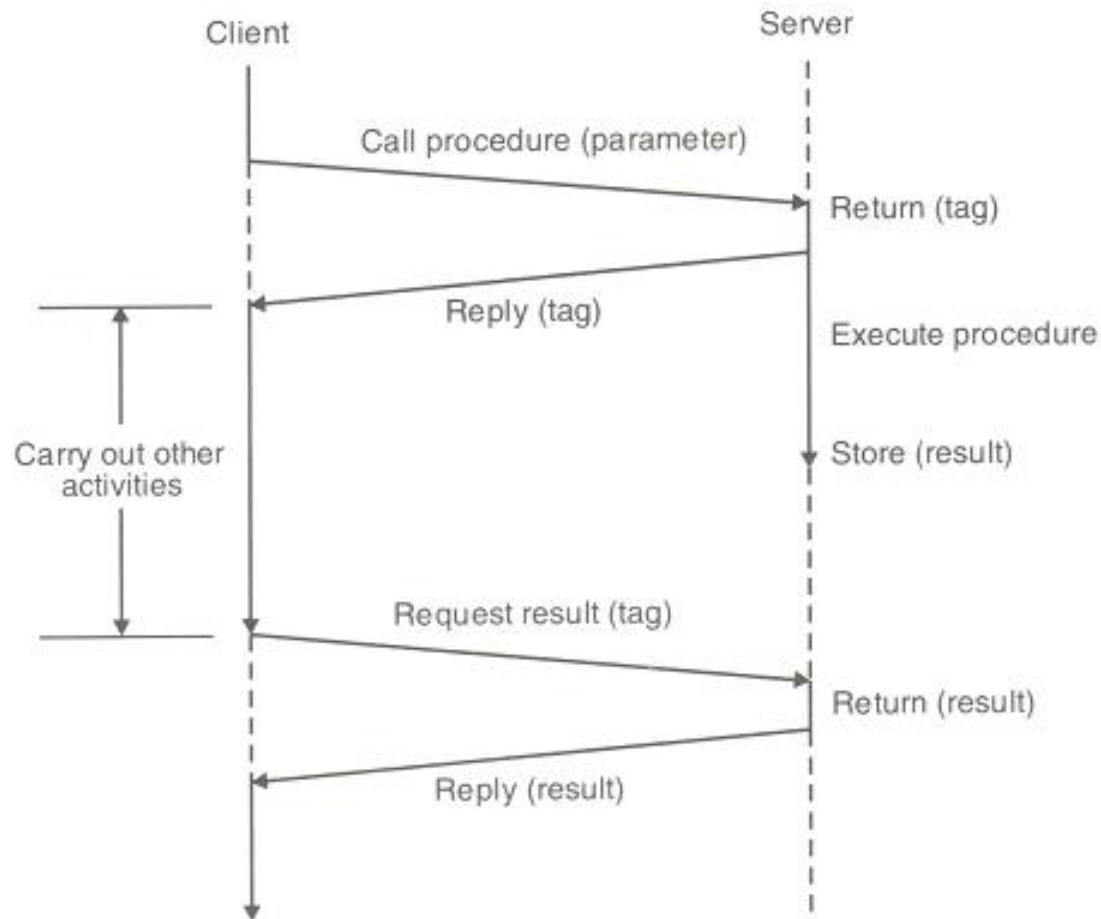
Optimizations for better Performance

- In Six Different Ways
 - Concurrent access to multiple servers
 - Serving multiple requests simultaneously
 - Reducing per-call workload of servers
 - Reply caching of idempotent remote procedures
 - Proper selection of timeout values
 - Proper design of RPC protocol specification

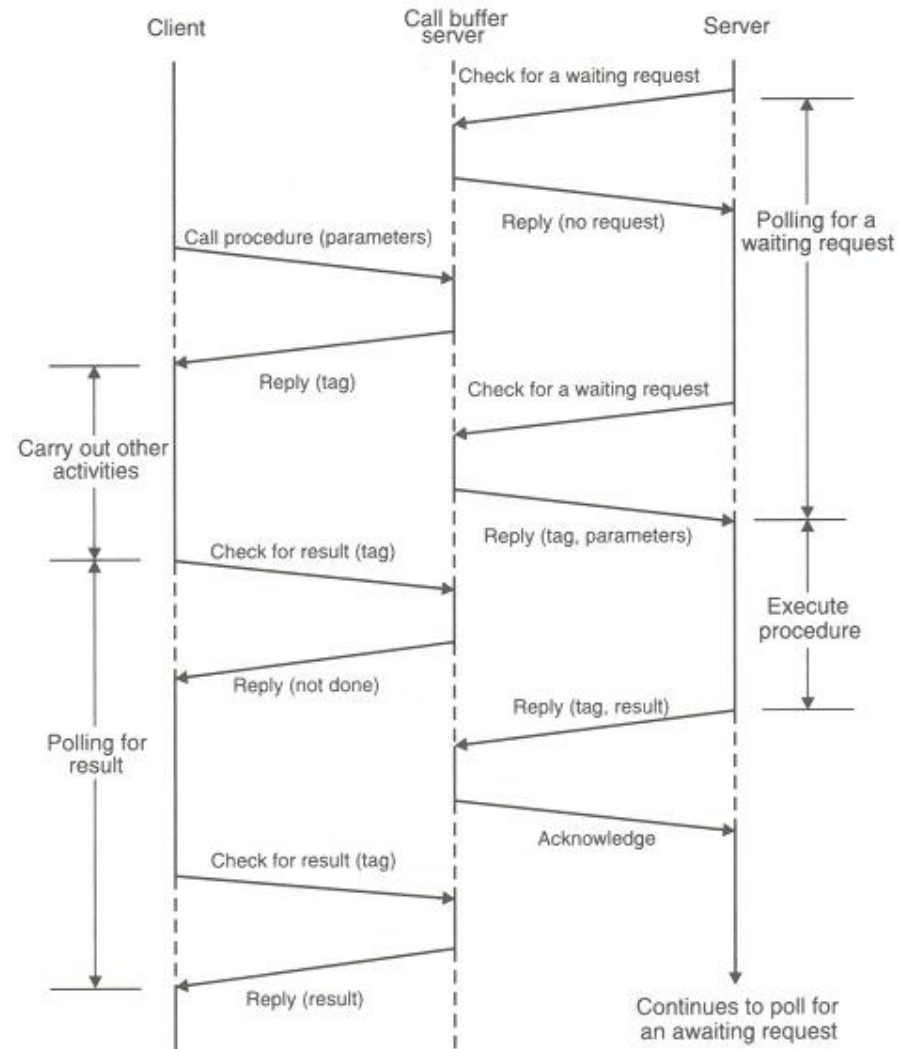
Concurrent Access to Multiple Server

- One of the following three may be adopted:
 - Threads
 - Use of Threads in the implementation of a client process where each thread can independently make remote procedure calls to different servers.
 - Addressing in underlying protocol should be rich enough to provide correct routing of responses.
 - Early reply approach [Wilbur and Bacarisse]
 - A call is split into two separate RPC calls- one passing parameters and other requesting the result
 - Server must hold the result causing congestion or unnecessary overhead.
 - Call buffering approach [Gimson]
 - Clients and servers do not interact directly but via a call buffer server
 - A variant of this approach was implemented in MIT (Mercury Communication System)

Early Reply Approach



Call Buffering Approach



Serving Multiple Requests Simultaneously

- Following types of delays are common-
 - A server, during the course of a call execution, may wait for a shared resource
 - A server calls a remote function that involves computation or transmission delays
- So the server may accept and process other requests while waiting to complete a request.
- Multiple-threaded server may be a solution.

Summary

- What is the purpose of IPC?
 - information sharing among two or more processes
- Differences between Synchronous and Asynchronous Communications?
 - When both the send and receive primitives of a communication between two processes use blocking semantics, the communication is said to be Synchronous; otherwise it is asynchronous
- List the Types of Failure in IPC
 - Loss of request message, Loss of response message, Unsuccessful execution of the request
- How to implement Idempotency?
 - Adding sequence number with the request message and Introduction of 'Reply cache'
- Three main types of Group Communications?
 - One to many, Many to one, Many to many
- One of the greatest challenges in Many to Many?
 - Ordered Delivery
- Name an all propose IPC protocol?
 - Remote Procedure Call (RPC)
- Name a few ways to optimise RPC?
 - Concurrent Access to Multiple Servers, Serving Multiple Requests Concurrently, Reducing Call Workload per Server
- Three different techniques for implementing Concurrent Access to Multiple Servers?
 - Threads, Early Reply, Call Buffering

References

1. Birman, K. P. and Renesse, R. V. *Reliable Distributed Computing with the ISIS Toolkit*. IEEE Computer Society Press, 1994.
2. Birrell, A. D. and Nelson, B. J. *Implementing remote procedure calls*. ACM Trans. Comput. Syst. 2(1), 39-59, 1984.
3. Wilbur, S. and Bacarisse, B. *Building distributed systems with remote call*, Software Engineering Journal, 2(5), 148-159, 1987.
4. R. Gimson. Call buffering service. Technical Report 19, Programming Research Group, Oxford University, Oxford University, Oxford, England, 1985.