



MONASH University

Information Technology

FIT3143 - LECTURE WEEK 8

PARALLEL ALGORITHM DESIGN -
PARTITIONING BASED ON MATRIX OPERATIONS

algorithm distributed systems **database**
systems **computation** knowledge ma
design e-business **model** data mining **int**
distributed systems **database** software
computation knowledge management **an**

ANANTH GRAMA • ANSHUL GUPTA
GEORGE KARYPIS • VIPIN KUMAR

Topic Overview

- Matrix Algorithms & Problem Statement
- Decomposition
- Decomposition – Fox's method

A portion of the content in the following slides were created by:

a) Gergel V.P., Nizhni Novgorod, **Introduction to Parallel Programming: Matrix Multiplication**, 2005.

b) Ananth Grama, Anshul Gupta, George Karypis, and Vipin Kumar, **“Introduction to Parallel Computing”, Addison Wesley, 2003.**

Introduction to Parallel Computing

Second Edition



Matrix Algorithms: Introduction

- Due to their regular structure, parallel computations involving matrices and vectors readily lend themselves to data-decomposition.
- Typical algorithms rely on input, output, or intermediate data decomposition.
- Most algorithms use one- and two-dimensional block, cyclic, and block-cyclic partitionings.

Problem Statement

Matrix multiplication:

$$C = A \cdot B$$

or

$$\begin{pmatrix} c_{0,0}, & c_{0,1}, & \dots, & c_{0,l-1} \\ & & \dots & \\ c_{m-1,0}, & c_{m-1,1}, & \dots, & c_{m-1,l-1} \end{pmatrix} = \begin{pmatrix} a_{0,0}, & a_{0,1}, & \dots, & a_{0,n-1} \\ & & \dots & \\ a_{m-1,0}, & a_{m-1,1}, & \dots, & a_{m-1,n-1} \end{pmatrix} \begin{pmatrix} b_{0,0}, & b_{0,1}, & \dots, & b_{0,l-1} \\ & & \dots & \\ b_{n-1,0}, & b_{n-1,1}, & \dots, & b_{n-1,l-1} \end{pmatrix}$$

The matrix multiplication problem can be reduced to the execution of $m \cdot l$ independent operations of matrix A rows and matrix B columns inner product calculation

$$c_{ij} = (a_i, b_j^T) = \sum_{k=0}^{n-1} a_{ik} \cdot b_{kj}, \quad 0 \leq i < m, \quad 0 \leq j < l$$


Data parallelism can be exploited to design parallel computations

Sequential Algorithm

```
// Sequential algorithm of matrix multiplication
double MatrixA[Size][Size];
double MatrixB[Size][Size];
double MatrixC[Size][Size];
int i,j,k;
...
for (i=0; i<Size; i++){
    for (j=0; j<Size; j++){
        MatrixC[i][j] = 0;
        for (k=0; k<Size; k++){
            MatrixC[i][j] = MatrixC[i][j] + MatrixA[i][k]*MatrixB[k][j];
        }
    }
}
```

Sequential Algorithm

- Algorithm performs the matrix **C** rows calculation sequentially
- At every iteration of the outer loop on **i** variable a single row of matrix **A** and all columns of matrix **B** are processed



- ***m·l*** inner products are calculated to perform the matrix multiplication
- The complexity of the matrix multiplication is $O(mnl)$.

Block-Striped Composition


- ***A fine-grained approach** – the basic subtask is calculation of one element of matrix **C***

$$c_{ij} = (a_i, b_j^T), \quad a_i = (a_{i0}, a_{i1}, \dots, a_{in-1}), \quad b_j^T = (b_{0j}, b_{1j}, \dots, b_{n-1j})^T$$

- *Number of basic subtasks is equal to **n^2** .*
- *As a rule, the number of available processors is less than **n^2** (**$p < n^2$**), so it will be necessary to perform the subtask scaling*

Block-Striped Decomposition

- ***The aggregated subtask*** – the calculation of one row of matrix **C** (the number of subtasks is n)
- ***Data distribution*** – rowwise block-striped decomposition for matrix **A** and columnwise block-striped decomposition for matrix **B**



Block-Striped Decomposition

Analysis of Information Dependencies

*Each subtask hold one row of matrix **A** and one column of matrix **B**,*

- At every iteration each subtask performs the inner product calculation of its row and column, as a result the corresponding element of matrix **C** is obtained*
- Then every subtask i , $0 \leq i < n$, transmits its column of matrix **B** for the subtask with the number $(i+1) \bmod n$.*

*After all algorithm iterations all the columns of matrix **B** were come within each subtask one after another*

Block-Striped Decomposition

Aggregating and Distributing the Subtasks among the Processors:

- *In case when the number of processors p is less than the number of basic subtasks n , calculations can be aggregated in such a way that each processor would execute several inner products of matrix A rows and matrix B columns. In this case after the completion of computation, each aggregated basic subtask determines several rows of the result matrix C .*
- *Under such conditions the initial matrix A is decomposed into p horizontal stripes and matrix B is decomposed into p vertical stripes.*
- *Subtasks distribution among the processors have to meet the requirements of effective representation of the ring structure of subtask information dependencies.*

Block-Striped Decomposition

Efficiency Analysis...


- *Speed-up and Efficiency generalized estimates*

$$S_p = \frac{n^3}{(n^3/p)} = p \qquad E_p = \frac{n^3}{p \cdot (n^3/p)} = 1$$

Developed method of parallel computations allows to achieve ideal speed-up and efficiency characteristics

Block-Striped Decomposition

- Another possible approach for the data distribution is the rowwise block-striped decomposition for matrices **A** and **B***



Block-Striped Decomposition


Analysis of Information Dependencies

- *Each subtask hold one row of matrix **A** and one row of matrix **B**,*
- *At every iteration the subtasks perform the element-to-element multiplications of the rows; as a result the row of partial results for matrix **C** is obtained,*
- *Then every subtask i , $0 \leq i < n$, transmits its row of matrix **B** for the subtask with the number $(i+1) \bmod n$.*

*After all algorithm iterations all rows of matrix **B** were come within every subtask one after another*

Block-Striped Decomposition – Fox's method

Data distribution – checkerboard scheme



Basic subtask is a procedure, that calculates all elements of one block of matrix **C**

$$\begin{pmatrix} A_{00} & A_{01} & \dots & A_{0q-1} \\ \dots & \dots & \dots & \dots \\ A_{q-10} & A_{q-11} & \dots & A_{q-1q-1} \end{pmatrix} \times \begin{pmatrix} B_{00} & B_{01} & \dots & B_{0q-1} \\ \dots & \dots & \dots & \dots \\ B_{q-10} & B_{q-11} & \dots & B_{q-1q-1} \end{pmatrix} = \begin{pmatrix} C_{00} & C_{01} & \dots & C_{0q-1} \\ \dots & \dots & \dots & \dots \\ c_{q-10} & C_{q-11} & \dots & C_{q-1q-1} \end{pmatrix}, \quad C_{ij} = \sum_{s=0}^{q-1} A_{is} B_{sj}$$

Block-Striped Decomposition – Fox's method

Analysis of Information Dependencies

- *Subtask with (i,j) number calculates the block \mathbf{C}_{ij} of the result matrix \mathbf{C} . As a result, the subtasks form the $q \times q$ two-dimensional grid,*
- *Each subtask holds 4 matrix blocks:*
 - *block \mathbf{C}_{ij} of the result matrix \mathbf{C} , which is calculated in the subtask,*
 - *block \mathbf{A}_{ij} of matrix \mathbf{A} , which was placed in the subtask before the calculation starts,*
 - *blocks \mathbf{A}_{ij}' and \mathbf{B}_{ij}' of matrix \mathbf{A} and matrix \mathbf{B} , that are received by the subtask during calculations.*

Block-Striped Decomposition – Fox's method

Analysis of Information Dependencies – during iteration l , $0 \leq l < q$, algorithm performs:

- The subtask (i,j) transmits its block A_{ij} of matrix A to all subtasks of the same horizontal row i of the grid; the j index, which determines the position of the subtask in the row, can be obtained using equation:

$$j = (i + l) \bmod q,$$

where \bmod operation is the procedure of calculating the remainder of integer-valued division,

- Every subtask performs the multiplication of received blocks A_{ij}' and B_{ij}' and adds the result to the block C_{ij}

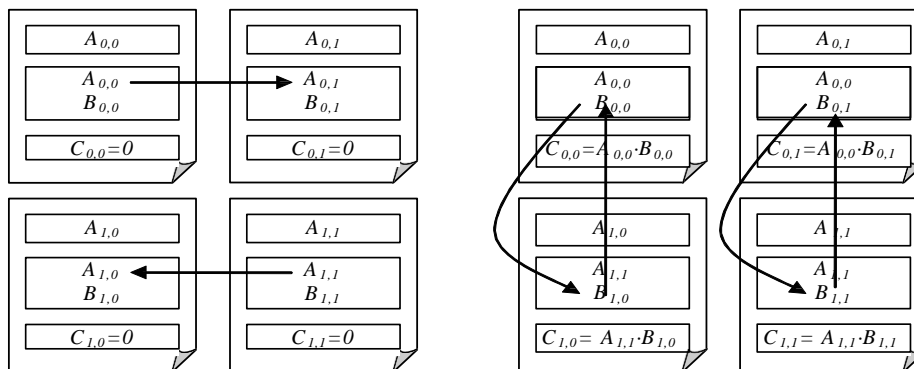
$$C_{ij} = C_{ij} + A_{ij}' \times B_{ij}'$$

- Every subtask (i,j) transmits its block B_{ij}' to the neighbor, which is previous in the same vertical line (the blocks of subtasks of the first row are transmitted to the subtasks of the last row of the grid).

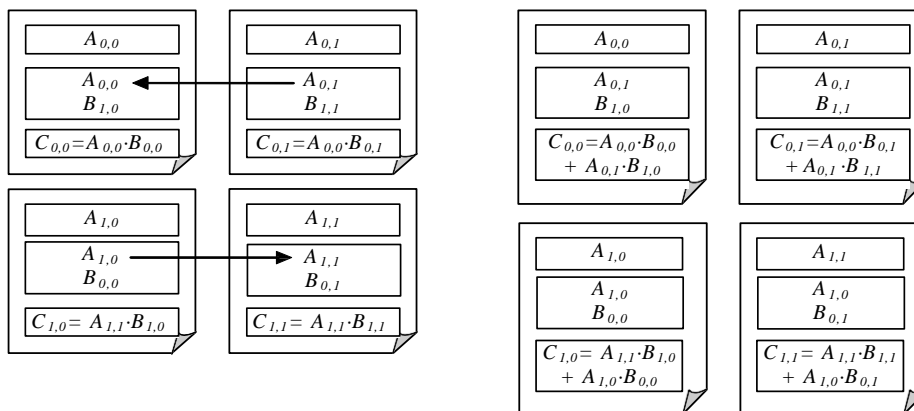
Block-Striped Decomposition – Fox's method

Scheme of Information Dependences

First Iteration



Second Iteration



Block-Striped Decomposition – Fox's method

Scaling and Distributing the Subtasks among the Processors

- *The sizes of the matrices blocks can be selected so that the number of subtasks will coincides the number of available processors p ,*
- *The most efficient execution of the parallel the Fox's algorithm can be provided when the communication network topology is a two-dimensional grid,*
- *In this case the subtasks can be distributed among the processors in a natural way: the subtask (i,j) has to be placed to the $p_{i,j}$ processor*

Block-Striped Decomposition – Fox's method

Scaling and Distributing the Subtasks among the Processors

- *The sizes of the matrices blocks can be selected so that the number of subtasks will coincide the number of available processors p ,*
- *The most efficient execution of the parallel the Fox's algorithm can be provided when the communication network topology is a two-dimensional grid,*
- *In this case the subtasks can be distributed among the processors in a natural way: the subtask (i,j) has to be placed to the $p_{i,j}$ processor*

In depth discussion & example

- Please refer to the enclosed report attached with these slides, “*Design and Implementation of Parallel Matrix Multiplication Algorithms using Message Passing Interface*” by Chin-Kit Ng for further in-depth discussion and code examples.
 - Serial matrix multiplication example
 - Bernstein analysis for data dependency
 - Parallel matrix multiplication examples using POSIX and MPI